

31338 Network Servers

32520 Systems Administration

Week 8

Networked filesystems: NFS & Samba

(Connect to Lab 8)

Professor Jian Zhang

School of Electrical & Data Engineering

University of Technology, Sydney

Network File System (NFS) Lab 8a

- Network File System (NFS)* is a distributed file system protocol originally developed by Sun Microsystems (Sun) in 1984, allowing a user on a **client** computer to **access** files which is much like a local storage and mounted locally **but** over a computer network
- From the **server** side, NFS is used to **export** local file system over NT
- NFS builds on the Open Network Computing Remote Procedure Call (ONC RPC) system.
- NFS is an open IETF standard defined in a Request for Comments (RFC), allowing **anyone** to implement the protocol. **On-going !**
- NFS is often used with Unix OS such as Solaris, AIX, HP-UX), Apple's macOS, Unix-like OS such as Linux and FreeBSD and also available to Mac OS, MS-DOS, Microsoft Windows etc.

IETF: Internet Engineering Task Force

NFS History

1. NFSv1 ('84): for **in-house** experimental purposes.
2. NFSv2 ('89): **UDP*** only, cannot guarantee delivery, 32-bit files
3. **NFSv3** (1995):
 - Support for **64-bit** file sizes to handle files larger than 2 gigabytes (GB);
 - **UDP-- Stateless, /TCP** – stateful**, support caching
 - Network Lock Manager (**NLM**) prevent data corruption in client sides
4. **NFSv4** (May 2015): 4.1, 4.2
 - Popular now, **TCP only**, support v3's feature
 - Stateful (TCP): stateless for v2 and v3
 - Consolidate multiple ports to **well-known TCP port 2049**
 - Built-in file-locking management feature (not NLM)
 - Cross-platform interoperability, including Windows
 - Improved security and strong authentication (Kerberos protocol)

**User Datagram Protocol (UDP)*

***Transmission Control Protocol (TCP)*

Stateless and Stateful NFS

- “Stateless” --
 - Simplicity
 - Makes recovery simpler if server crashes
 - **Client** must keep info about open files & current location in file
 - **Client** pings back to check state of file
 - No consistency guarantee
- “Stateful” --
 - all state information is stored on **both the client and the server** when they are active and recovered mutually in the event of an outage.

Open Network Computing Remote Procedure Call (ONC RPC) system

- NFS builds on ONC RPC
 - **RPC** is a programming mechanism to allow one application to transparently invoke procedure/method calls on a remote machine (network call)
 - Many varieties of RPC – standard UNIX one is "Sun RPC"
 - NFS, **NIS*** and a few other things built on Sun RPC
 - Independent from machine types, operating systems, and network architectures
- Need to run a **daemon** that maps RPC addresses into the port numbers that the service is listening on
 - The daemon is called: **rpcbind**
 - previously known as **portmap**

* Network Information Service (NIS) is a distributed database that allows you to maintain consistent configuration files throughout your network – looks like a yellow page

Typical NFS Implementation

- Assuming a Unix-style scenario in which **one machine (the client) needs access to data stored on another machine** (the NFS server):
 1. The **server implements NFS daemon processes**, running by default as `nfsd`, to make its data generically available to clients.
 2. The server administrator determines what **to make available**, exporting the names and parameters of directories, typically using the `/etc/exports` configuration file and the **`exportfs` command**.
 3. The **server** security-administration ensures that it can recognize and approve **validated clients**.
 4. The **server** network configuration ensures that appropriate **clients can negotiate** with it through any firewall system.
 5. The **client** machine requests **access to exported data**, typically by **issuing a mount** command. (The client asks the server (`rpcbind`) which port the NFS server is using, the **client** connects to the NFS server (`nfsd`), `nfsd` passes the request to `mountd`)
 6. If all goes well, users on the client machine can then view and interact with mounted **filesystems on the server** within the parameters permitted.
- **Note that automation of the NFS mounting process using `/etc/fstab` and/or automounting facilities**

NFS Server (Lab 8a)

- Packages: **nfs-servre** **nfs-utils**, **rpcbind** & **nfs4-acl-tools**
- Configure file for which directories to export and how
 - **/etc/exports** (on machine **'ens37'** 10.0.2.1/24 with its subnet)

Directory to share

client machines that will have access to the directory

/share/IT_Projects client1(ro) # Lab 8a

/opt/perl client1(rw,no_root_squash) client2(ro)

/opt/general 192.168.3.0/24(rw,root_squash) # allow 192.168.3.0 - 255

Option Definitions

rw – allows both read and write access on the file system.

ro = allow clients read only access to the share

sync – tells the NFS server to write changes to disk before reply (applies by default).

all_squash – maps all UIDs and GIDs from client requests to the anonymous user.

no_all_squash – used to map all UIDs and GIDs from client requests to identical UIDs and GIDs on the NFS server.

root_squash – maps requests from root user or UID/GID 0 from the client to the anonymous UID/GID.

- Turn on (CMD) – *systemctl start rpcbind and systemctl start nfs server*

NFS Server

- Start the NFS server daemons: **nfsd, mountd:**

systemctl start nfs-server

The other services that are required for running an NFS server or mounting NFS shares such as rpcbind, **rpc.mountd**, lockd, will be automatically started.

- The configuration files for the NFS server are:
 - /etc/**nfs.conf** – main configuration file for the NFS daemons and tools.
 - /etc/**nfsmount.conf** – an NFS mount configuration file.
- **rpc.mountd**
 - When an NFS client requests a mount, **rpc.mountd** checks **/var/lib/nfs/etab** to verify:
 - which directories are exported,
 - which clients are allowed,
 - and with what options (rw, ro, root_squash, etc.). The workflow is
- To export system, run the **exportfs** command

-a means export or unexport all directories,

-r means reexport all directories

-v enables verbose output.

- Display current export list: **exportfs -s** or **-v**

Examples: exportfs -a; exportfs /home/notes; More details*

* https://sites.ualberta.ca/dept/chemeng/AIX-43/share/man/info/C/a_doc_lib/cmds/aixcmds2/exportfs.htm

Enable NFS in Firewall

- Enable NFS services in firewall
 - GUI: **firewall-config**: enable NFS and make it permanent
 - CMD:
 - **firewall-cmd --permanent --add-service=nfs**
 - **firewall-cmd --permanent --add-service=rpc-bind**
 - **firewall-cmd --permanent --add-service=mountd**
 - **firewall-cmd --reload**

NFS Client (Lab 8a)

- On the client node(s), install the necessary packages to access NFS shares on the client systems.

dnf install nfs-utils

- Clients mount a filesystem that is located on a server

showmount -e 10.0.2.1

- Create a local file system/directory (yellow page) for mounting the remote NFS file system and mount it as a ntf file system.

mkdir -p /media/perl

mount -t nfs4 10.0.2.1:/opt/perl /media/perl

In the directory /media/perl*, Files are written by perl*

* Perl is a programming language (for text manipulation) that can be used to perform tasks that would be difficult or cumbersome on the command line.

NFS Client (recall lecture/lab 7!)

- To enable the mount to **persistent** even after a system reboot:
- Add entry to **/etc/fstab** to access remote directory (on machine 'client1'):

#	Device	Mount point	FS type	FS options	F	P
	server1:/opt/perl	/media/perl	nfs	rw	0	0

The last two digits: dump and pass

dump: A number indicating whether and how often the file system should be backed up by the dump program. 0 means never automatically backed up.

Pass: A number indicating the order in which the fsck program will check the devices for errors at boot time. 1: root file system; 2: check after root; 0: not check for all other devices

rw: read and write; ro: read only

e.g., noauto, nosuid, rw can be the options field (namely "rw" field in above example)

- Then mount:

```
mount /media/perl
```

Automount

- NFS becomes complex in a system with several servers and many clients
 - Lots of copies of `/etc/fstab` to maintain
 - Network stops if server **crashes**
- Automount is more flexible and easier to manage
 - Install `autofs` service and edit `/etc/auto.master`
- Clients mount directories when they need them
 - directories unmounted when **no longer used**
 - directories may be available from **alternative servers**

- Which of the following /etc/exports directives would be important for trying to create a read-only share's configuration record? (Choose all that apply.)
 - A. ro
 - B. rw
 - C. async
 - D. sync

SMB and Samba (Lab 8b)

- There is no **user authentication** when we configure the NFS. Hence,
- SMB = Server Message Block: is a **client-server communication protocol** used for sharing **access to files, printers, serial ports and other resources on a network.**
- Samba: **a** software implements of SMB
 - Samba allows files to be shared **across windows and Linux and printers** and other devices in a simple and seamless manner.
 - Samba is an **open-source** SMB implementation
 - On Windows: Common Internet File System (CIFS): A software implemented SMB protocol, native network filesystem protocol for Windows

Samba Server (Win/Linux)

- On Windows: Right click to share – easy job!
- On Linux:
 - Package: samba using `dnf install samba` to install
 - Service/daemon: `smbd` and `nmbd`
 - `smb (d)`: samba server, provide share locking, user authentication
 - `nmb (d)`: **NetBIOS*** over TCP/IP name service, like RPC in NFS
 - Designed by IBM, only support LAN initially, now both LAN and WAN
 - `systemctl start/restart/enable smb (nmb)`
 - Log file: `/var/log/samba/*`
 - Configuration file (see the next slide!):
`/etc/samba/smb.conf`

*** NetBIOS (Network Basic Input/Output System) is a network service that enables applications on different computers to communicate with each other across a local area network (LAN).**

Samba Server (Win/Linux)

- Configuration: **global section** in **/etc/samba/smb.conf**

[global] – (more info see *)

workgroup = WORKGROUP # This is default workgroup for Windows machines

security = user # together with next line, enables Linux system users to log in to the Samba server.

passdb backend = tdbsam # backends" (or ways) of storing login access with Samba.

printing = cups # Common UNIX Printing System to share printers via Samba host systems

printcap name = cups # users connects to any printer specified in the local host's printcap file

load printers = yes # Enable the automatic printer sharing

cups options = raw # File is print-ready, do not process any further

netbios name = MYSAMBASERVER # used to specify a server name that is not tied to the hostname, this becomes the machines's "Samba hostname"

interfaces = 10.0.2.0/24 127.0.0.0/8 #used to configure Samba to listen on multiple network interfaces. (two subnets!)

hosts allow = 10.0.2. # the hosts allowed to connect (Local subnet).

Samba Server Lab 8b

- Configure “Global section” in `/etc/samba/smb.conf`
 - Highly commented (see `smb.conf.example`), read and uncomment as required

Copy the settings from example file to [global] section of `smb.conf` [global]

```
workgroup = MYGROUP
server string = Samba Server Version %v
;
netbios name = MYSERVER
;
interfaces = lo eth0 192.168.12.2/24 192.168.13.2/24
;
hosts allow = 127. 192.168.12. 192.168.13.
```

Then modify the settings based on the task 1 of Lab 8b

- `workgroup` (e.g. set to “WORKGROUP”, which is the default for Windows machines)
- `netbios name` (set to “MYSAMBASERVER” – this becomes the machine's “Samba hostname”)
- `interfaces` (e.g. use “10.0.2.0/24” and “127.0.0.0/8”)
- `hosts allow` (set to “10.0.2.” – note the unusual usage of dots)

Samba Server (Lab 8b)

- Configure "shares" in `/etc/samba/smb.conf`

- Highly recommended (see `smb.conf.example`), read and uncomment as required

- `[homes]` *# [homes] is a special entry – notice it has no path attribute*

- `comment = Home Directories`

- `valid user = %s, %D%w%S`

- `browseable = Yes`

- `read only = Yes`

- `Inherit acls = Yes # Access Control Lists`

- Note: `[homes]`, `[opt]` & `[public]` are all share entries

- Log and troubleshooting

- syntax check: `testparm /etc/samba/smb.conf`

- Log file: `/var/log/samba/*`

- Firewall

- `firewall-config` (GUI!): public runtime & permanent samba service

`[global]` = server-wide settings (authentication, logging, networking).

`[share]` = per-folder configuration (who can access, read/write rules, permissions).

Lab 8b: Tasks 4 and 5

`[opt]` *# general entry*

`path=/opt`

`public=yes`

`writable=no`

`browseable = yes`

`[public]`

`comment = Shared public
directory`

`path = /pub`

`writable = yes`

`public = yes`

Samba security

- As with NFS, basic Samba access control comes from configuring options on shares, e.g.
 - **browseable** (can browse directory)
 - **writable** (can write to directory)
 - **public** (publically accessible)
 - **Read only** (publically read only)

Instead of 'writable = Yes' you can also say 'read only = Yes'

- Samba configuration can get quite complicated
 - Not too bad for simple sharing
 - For serious use, read the long, long manual

Samba Configuration (Lab 8b)

- Samba security is **user-based**
 - users require username & password to access Samba shares
 - password is encrypted during transmission, file data is not
- Samba maintains its **own authentication database**
 - **Separate from UNIX passwd file**
 - **Tools to keep Samba/UNIX passwords in sync**
 - or have Linux authenticate against Samba database (PAM)
 - or authenticate against central Windows PDC or LDAP server
 - **pdbedit** CMD to create **new users** and edit Samba password database
 - **pdbedit -a -u newuser #** -a: to add a user into the database. It needs a user name specified with the -u switch. When adding a new user, pdbedit will also ask for the password to be used.
 - **pdbedit -x -u olduser #** -x: to delete an account from the database. It needs a username specified with the -u switch.
 - **pdbedit -L #** lists all the user accounts present in the users database.

Note: Don't forget to run “systemctl restart smb” and “systemctl restart nmb*” & enable them!

Samba Client (Win/Linux) Lab 8b

- Task 2: Testing Samba from **Linux machines**

- Permanent access in `/etc/fstab`

```
//server/public /pub cifs username=xxx,password=xxx 0 0
```

- Temporary:

- Connect: `smbclient -U peter //10.0.2.1/peter # user login`

```
smbclient -L 10.0.2.1 # list=HOST
```

- Check files: `ls, dir`

- Download: `get mywinfile.txt /tmp/win`

- Upload: `put /etc/samba/smb.conf samba.txt`

- normal directory: `\\10.0.2.1\shared` #create it, lab task 4

- Task 3: Testing Samba from **Windows machines**

- if kernel has smbfs /cifs support

- Access by server's IP address or NetBIOS name

- File Explore  in the address bar type `\\10.0.2.1\peter`

- normal directory: `\\10.0.2.2\shared` #create it the lab task 5

Q & A

- Which of the following Samba directives can be used to set **whether passwords are required to access this file share service?** (Choose the best answer.)
 - A. public** → synonym for guest ok. If set to yes, the share allows guest access **without requiring a password**.
 - B. browseable** → only controls whether the share is visible in "Network Neighborhood / Explorer". Doesn't affect password.
 - C. guest only** → means users can only connect as guests (never with a password).
 - D. writable** → controls write access, not password requirement.
 - E. read only** → controls read/write, not password requirement.
- After class practice: how to share directories with different access control for different groups?
 - User, Group
 - File permission, ownership
 - Samba configuration: group? **man smb.conf**
 - Samba doesn't have its own group concept; it relies on **Unix/Linux groups**. Instead using: **admin users, valid users, create mask, directory mask**